In the following assignments, we want you to learn how to use an Azure SQL database from scratch. Your first step will thus be to set up a database in the Azure service and import your data. This step may seem tedious but it is crucially important. We want you to be able to continue using Azure after the class ends. For this, you need to know how to use the system starting from nothing.

NOTE: These steps will take some time to complete, so start early!

A. Setting up an Azure SQL Database

Step 1: Create an Azure account and login to Azure portal

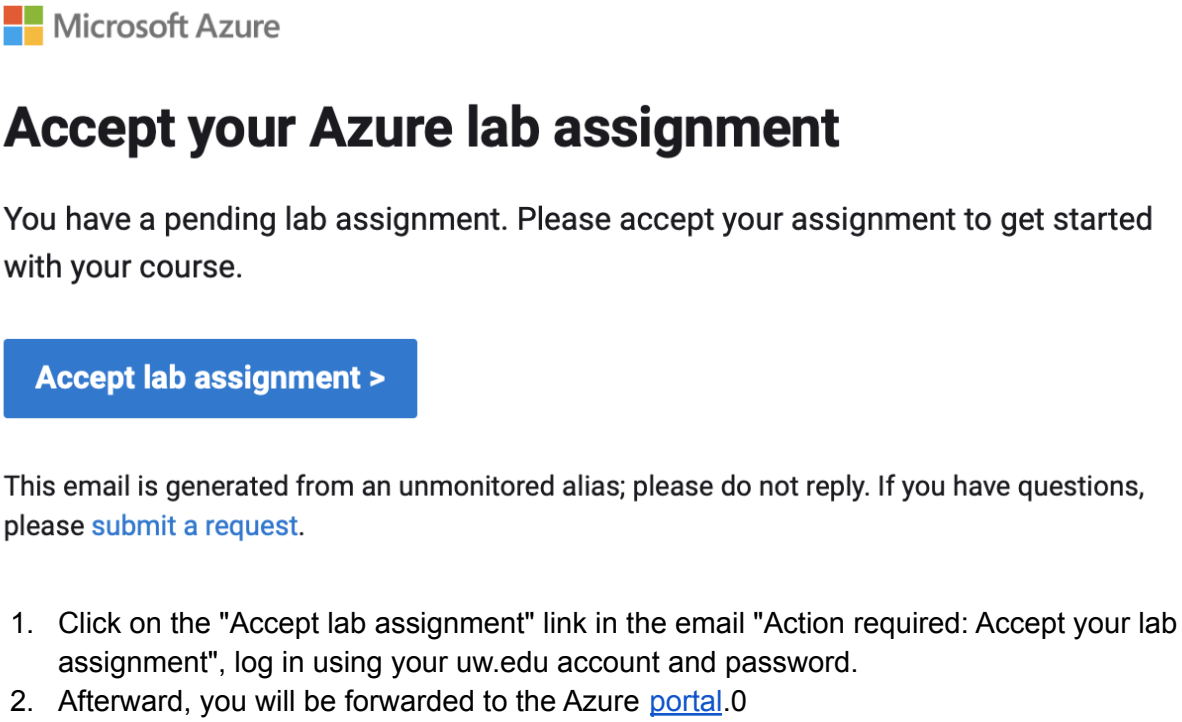

1. Click on the "Accept lab assignment" link in the email "Action required: Accept your lab assignment", log in using your uw.edu account and password.
2. Afterward, you will be forwarded to the Azure portal.0

3. Under “Action needed”, click “Setup lab”. This might take a while.

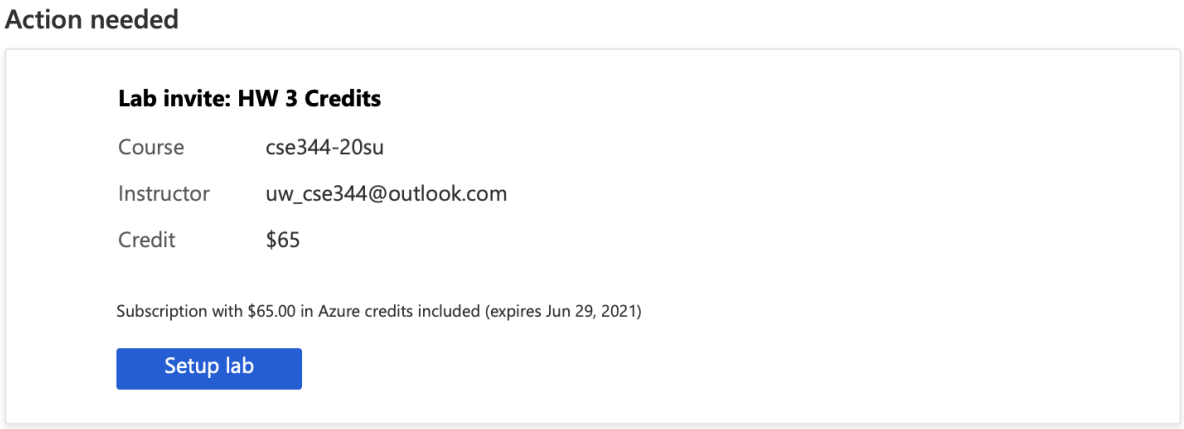

NOTE: The above screenshot is showing an old invite for a different class and your information will be different. If you have any questions, please contact the course staff.

Step 2: Learn about Azure SQL Server

Spend some time clicking around, reading documentation, watching tutorials, and generally familiarizing yourself with Azure and SQL Server.
Resources:

- SQL Server Tutorial For Beginners
- SQL Server Technical Documentation
- SQL Server Tutorial
- Transact-SQL Reference

Step 3: Create a database

1. Navigate to the Azure portal.

2. Select “Create a resource” under Azure services.

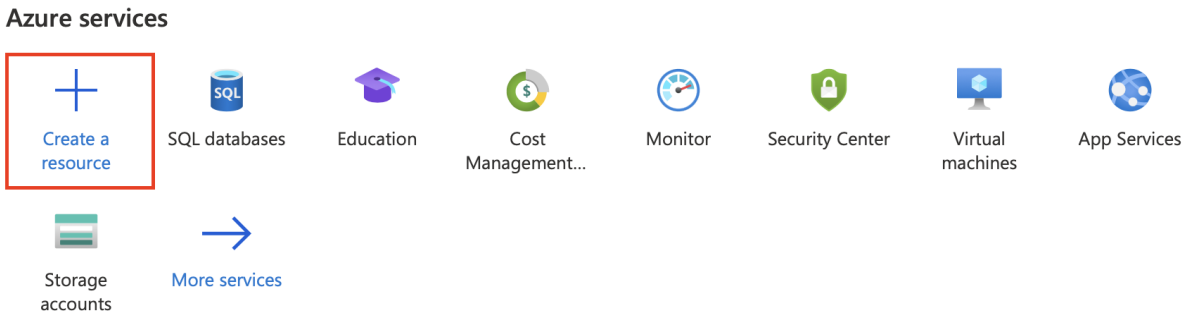

3. Select “Databases” under Azure Marketplace, and then select “SQL Database” under Featured.

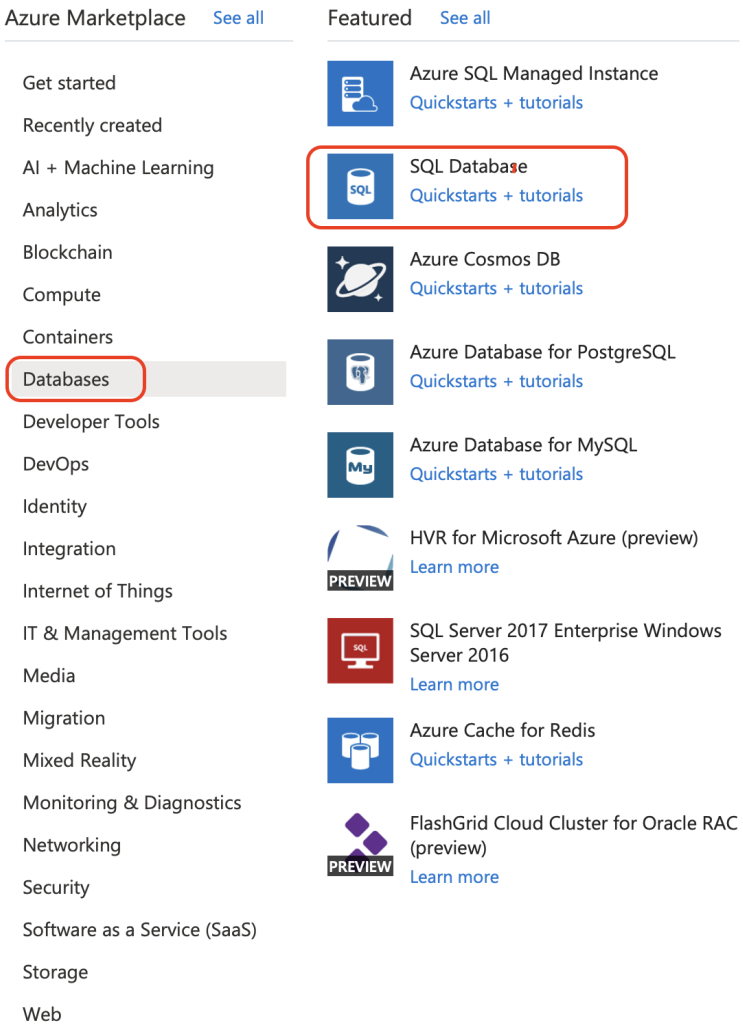

4. Make sure that your “Subscription” is **NOT** set to “Free Trial” or “Azure for Students”. The correct subscription name will either:
    a. mention “Homework 3”
    b. or mention “Microsoft Azure Sponsorship”

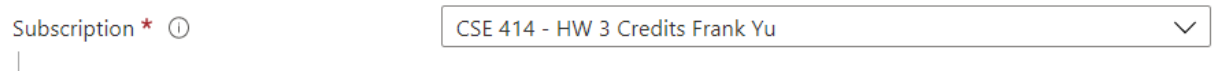

5. Under Subscription, create a new resource group with a name.

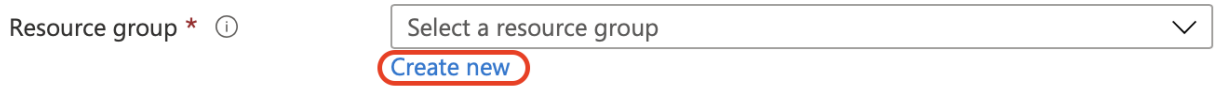

6. Choose a database name.

7. Create a new server. Enter your server name, admin login, and password. Choose the server location based on where you are at. For example, if you’re located in Seattle, choose “West US 3”. Make sure Authentication method is set to 'Use SQL authentication'.

8. Make sure that “Want to use SQL elastic pool?” is set to **NO.**

9. Next, for “Compute + storage”, click on “Configure database”, a second panel will open.

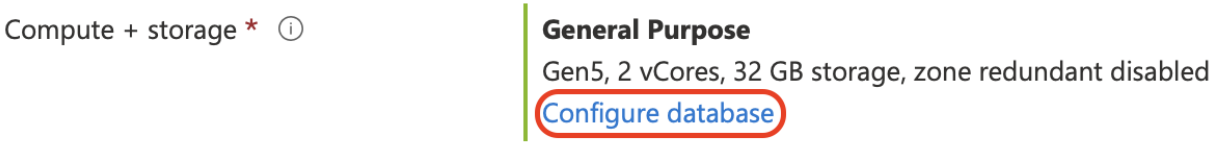

10. On the banner, click on “Looking for basic, standard, premium?”.

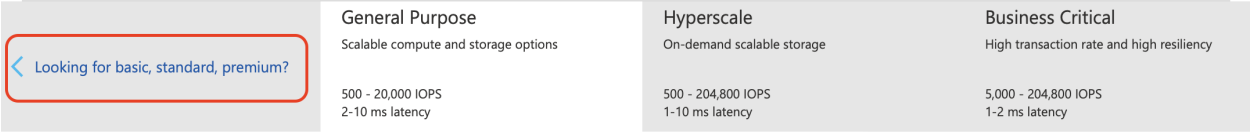

11. Select “Standard”.

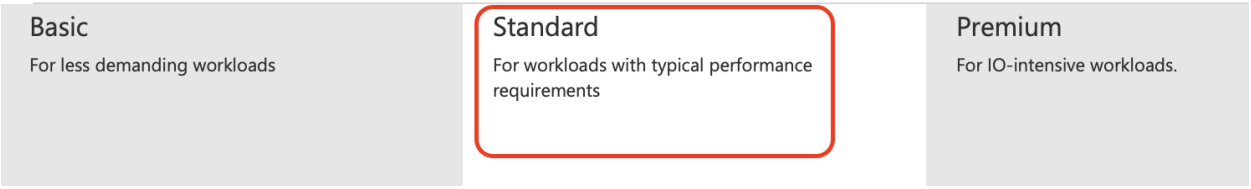

12. Check your DTUs, Data max size, and estimated cost/month.

13. Click on “Apply” at the bottom of the page.

14. Click on “Next: Networking” at the bottom of the page.

15. For “Connectivity method”, select “Public endpoint”.

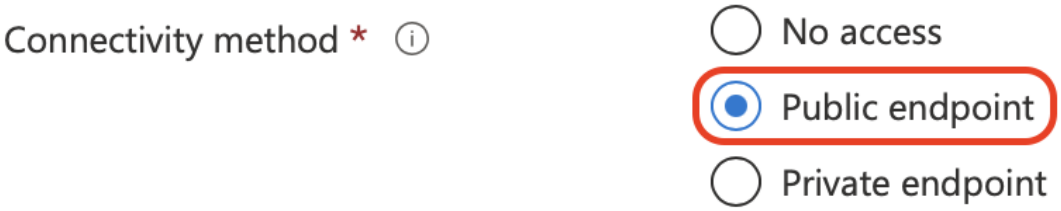

16. Under “Firewall rules”, make sure that “Allow Azure services and resources to access this server” is set to “Yes”.

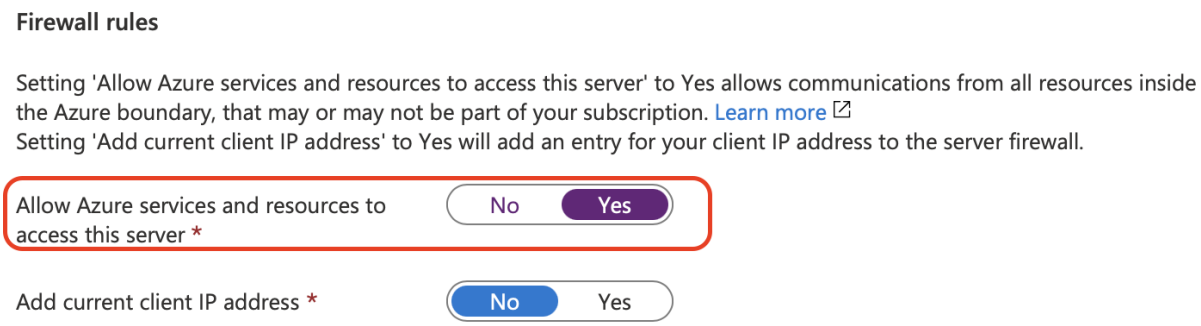

17. Click on “Next: Security” at the bottom of the page.

18. Make sure that “Enable Azure Defender for SQL” is set to “Not now”.

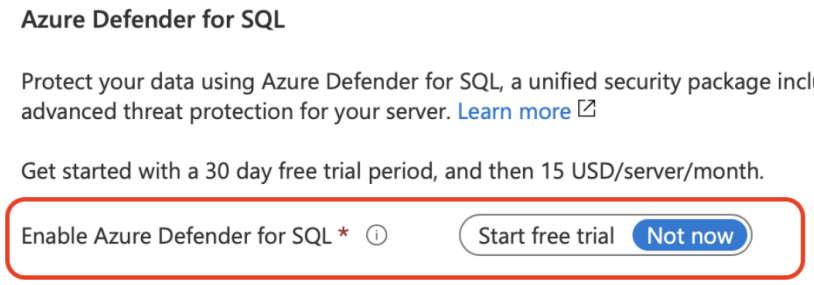

19. Click on “Next: Additional settings” at the bottom of the page.

20. Make sure that “Use existing data” is set to “None”.

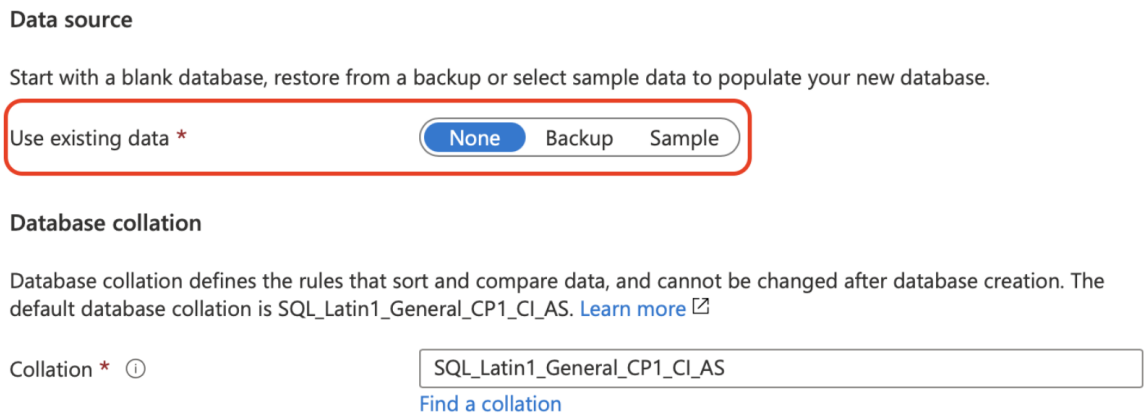

21. Click on “Review + create” at the bottom of the page.

22. Click on “Create” at the bottom of the page.

23. Wait a few minutes! If you encounter any issues, please contact the course staff!

24. Once the deployment is complete, navigate to Azure portal. Under “Azure services”, click on “SQL database”.

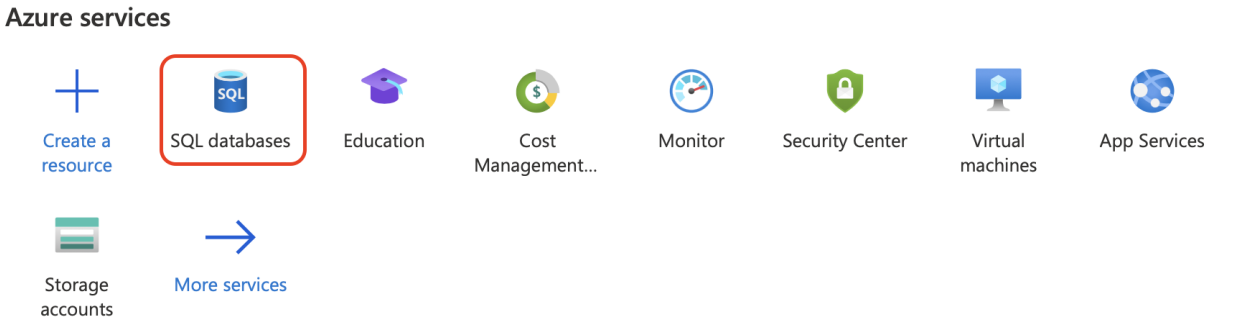

25. Click on the database that you just created.

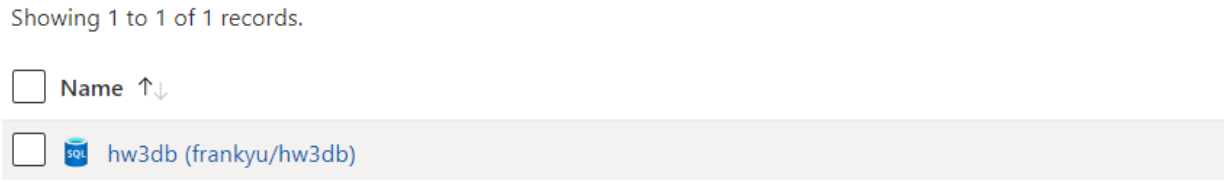

26. At the top of the page, click on “Set server firewall”.

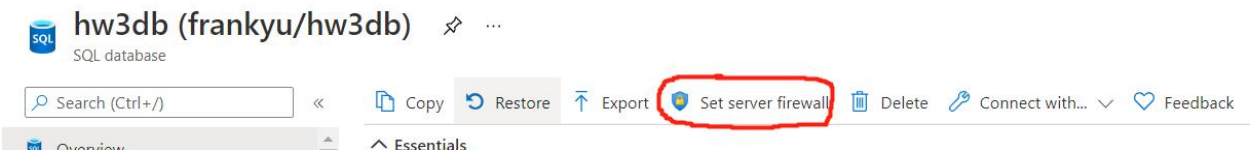

27. Add a new rule that allows connections from any client. For “Start IP”, type in “0.0.0.0”, for “End IP”, type in “255.255.255.255”. Remember to save your settings!

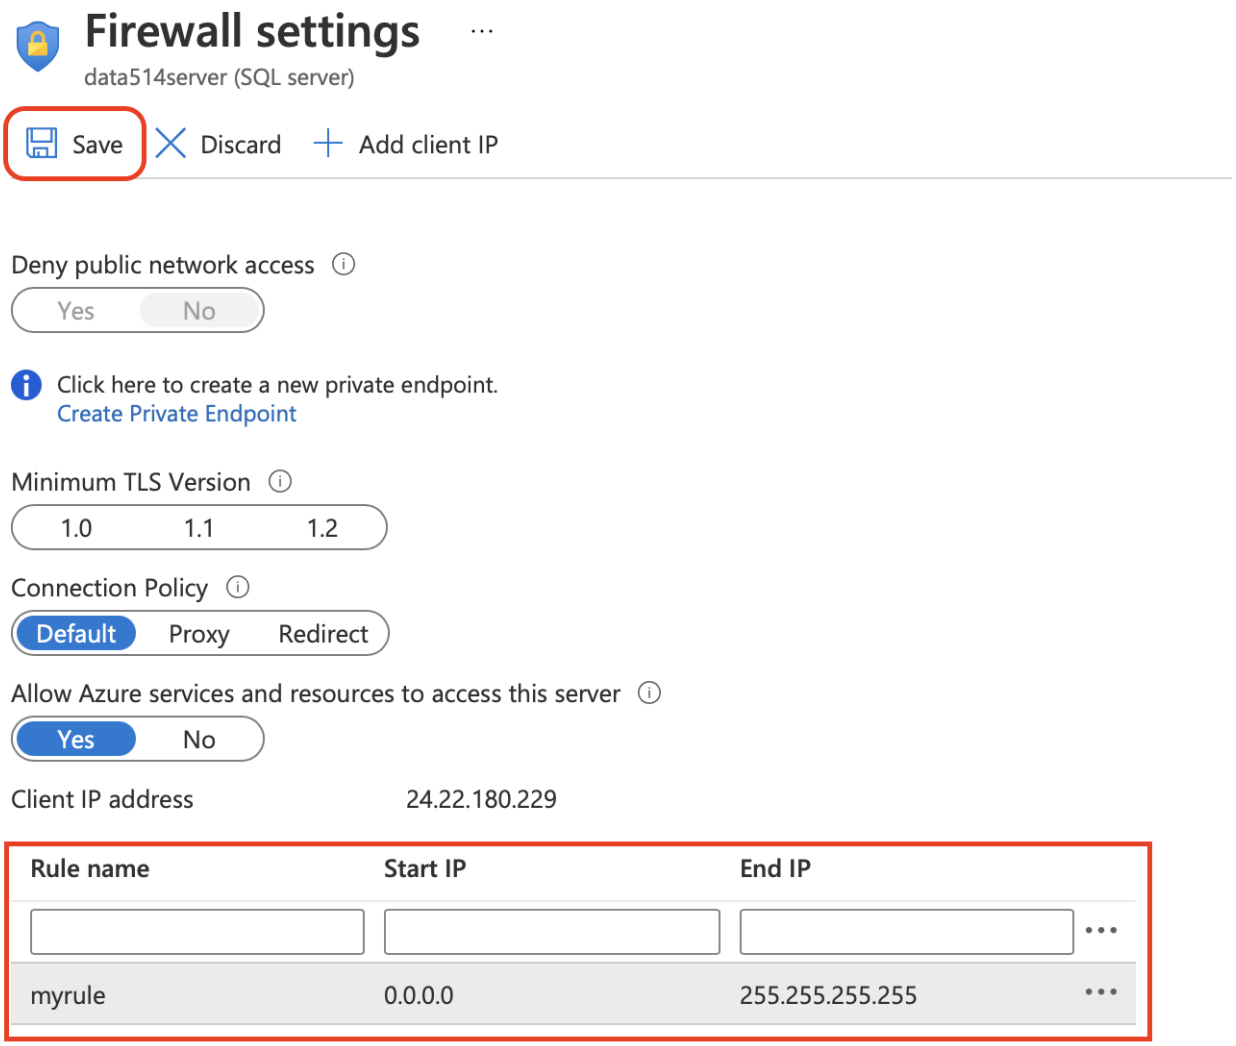

28. You have successfully set up your first SQL database on Azure!

Step 4: Try out the built-in Query editor

The simplest way to play with the database is using the built-in Query editor in the Azure portal. To launch this:

1. Navigate to the SQL database you just created (see 22 & 23 in the above section).

2. Click on “Query editor (preview).

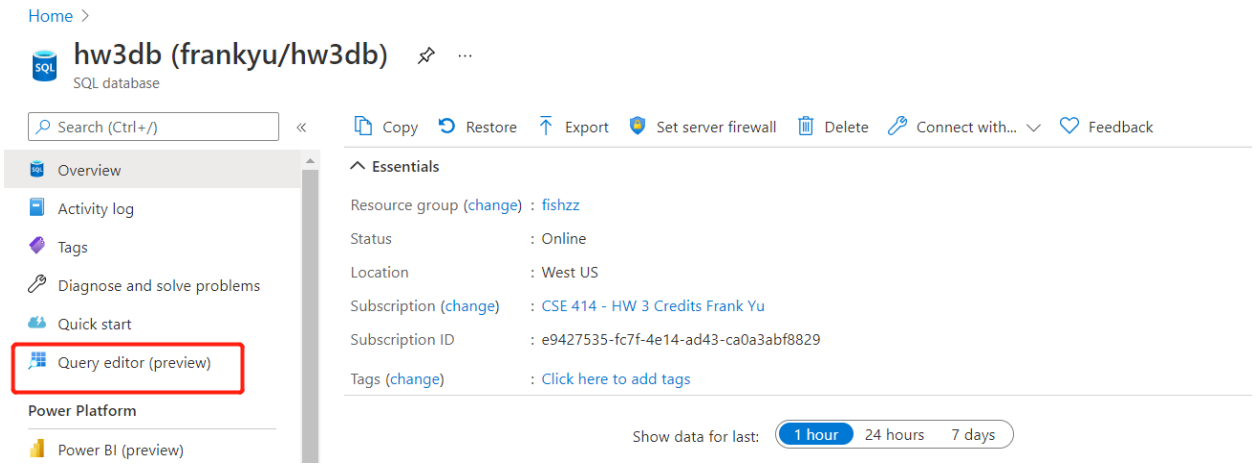

3. Log in using the username and password you created during step 7 in the above section.

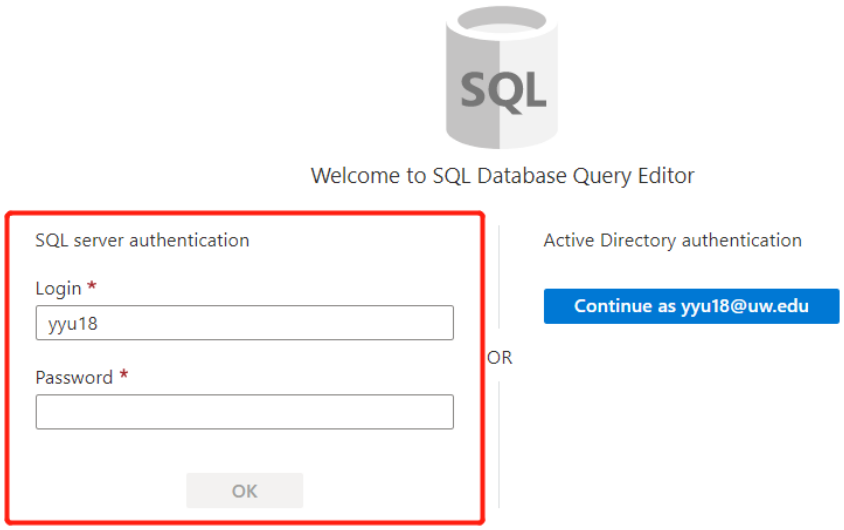

4. That’s it!

Step 4.1: Use an IDE! (Optional)

Besides using the built-in query editor, another option would be to connect VS Code to your database. Visual Studio Code is a freeware source-code editor made by Microsoft for Windows, Linux, and macOS.
Read the instructions here to set up your VS Code! More specifically, follow the steps in “Install and Start Visual Studio Code”, “Install the mssql extension”, “Create or open a SQL file”, and most importantly, “Connect to SQL Server”.
When you’re connecting to SQL Server, you will be prompted to enter your Server name. This can be found in the dashboard for your SQL database.

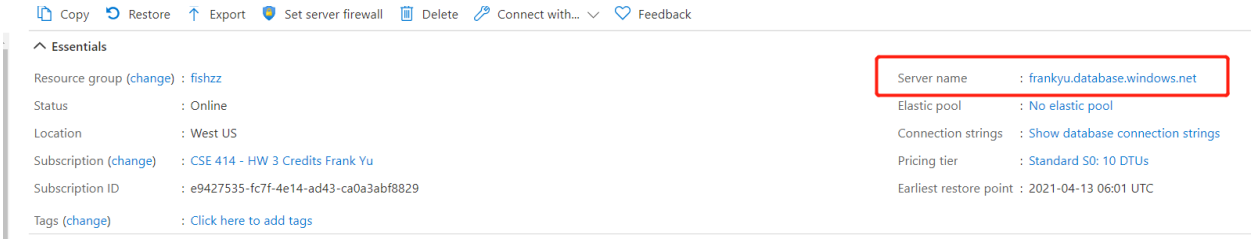

Step 5: Ingest the data!

Important: You would need to create your tables first before ingesting the data!
Follow the following steps to ingest the data from our remote host to your SQL database.

1. Create an external data source. Copy and paste the following code in your Query editor and execute.

```
CREATE EXTERNAL DATA SOURCE cse414blob
WITH (TYPE = BLOB_STORAGE,
LOCATION = 'https://cse344.blob.core.windows.net/cse344-22wi'
);
```

2. Execute your create table statements; you can just copy over your `CREATE TABLE` statements from hw2 (don’t copy the `` `PRAGMA foreign_keys=ON` `` or the ‘.mode’ statements; those are specific to SQLite).

3. Copy and paste the following code in your Query editor and execute.

```
bulk insert Carriers from 'carriers.csv'
with (ROWTERMINATOR = '0x0a',
DATA_SOURCE = 'cse414blob', FORMAT='CSV', CODEPAGE = 65001, --UTF-8
encoding
FIRSTROW=1,TABLOCK);
```

```
bulk insert Months from 'months.csv'
with (ROWTERMINATOR = '0x0a',
DATA_SOURCE = 'cse414blob', FORMAT='CSV', CODEPAGE = 65001, --UTF-8
encoding
FIRSTROW=1,TABLOCK);

bulk insert Weekdays from 'weekdays.csv'
with (ROWTERMINATOR = '0x0a',
DATA_SOURCE = 'cse414blob', FORMAT='CSV', CODEPAGE = 65001, --UTF-8
encoding
FIRSTROW=1,TABLOCK);

-- Import for the large Flights table
-- This last import might take a little under 10 minutes on the provided
server settings

bulk insert Flights from 'flights-small.csv'
with (ROWTERMINATOR = '0x0a',
DATA_SOURCE = 'cse414blob', FORMAT='CSV', CODEPAGE = 65001, --UTF-8
encoding
FIRSTROW=1,TABLOCK);

-- Indexes to make your query run faster (optional, but recommended)
-- This will take about 3-4 minutes

create index Flights_idx1 on Flights(origin_city,dest_city,actual_time);
create index Flights_idx2 on Flights(actual_time);
create index Flights_idx3 on Flights(dest_city,origin_city,actual_time);
```

4. In order to verify that your imports were successful, do some SELECT COUNT(*) statements. The correct input should have:
    - 1594 rows for Carriers.
    - 12 rows for Months.
    - 8 rows for Weekdays.
    - 1148675 rows for Flights.